

Feature Selector and Tracker

Assignment 2

Gabriele Padovani
Artificial Intelligence Systems
Trento, Italy
gabriele.padovani@studenti.unitn.it

I. INTRODUCTION

To offer a quantitative estimate of the performance of the extractors used, as well as of the tracking, the average frames per second achieved during execution is calculated. This metric, though not perfect gives an idea of how performant an extractor is compared to the others. Since the number of features found by different algorithms may change, a secondary metric is calculated, keeping track of how many points are returned for each frame.

It is important to stress that a higher number of points found does not imply that the feature extractor is better, as, for example with FAST, tens of thousands of points returned render real-time tracking infeasible.

II. FEATURE SELECTOR APPROACHES USED

The feature extractors evaluated for this second Computer Vision Assignment are:

- **SIFT**: applies an increasingly intense Gaussian filter to the reference image, and looks for maximum and minimum value points, after performing pixel-wise difference.
- **FAST**: corner detector, compares pixel values in a specified radius, and is designed to offer high performance;
- **ORB**: similar to FAST, but also takes into account the rotation of features;
- **Good Features to Track**: compares the diversity matrices of neighboring pixels, looking for points with strong eigenvalue, both in the x and y components.

III. TRACKERS USED

To offer tracking capability, Lukas Kanade Optical Flow was used in conjunction with Good features to track, while SIFT and ORB were both tested with OpenCV's Brute Force matcher and Flann-based matcher.

The former works by refining an initial estimate of a displacement vector, taken from the image at a very low resolution, and corrects it by checking it against the same image at increasingly higher definitions.

On the other hand, the brute force matcher works by comparing, each feature in the first image, all the points in the second one, and returning the best match. The Flann-based matcher promises instead higher performance on large sets of features, by learning on a set of descriptors, and utilizing a more efficient type of distance metric, to quickly exclude false matches.

IV. PERFORMANCE COMPARISON

As explained in the introductory section, two metrics were calculated, to offer some kind of quantitative evaluation of these algorithms:

Feature extractors (Alone)

	Features per Frame	Average FPS
FAST	23323.17	5.43
GFTT	25.0	4.26
ORB	1000.0	5.72
SIFT	3339.3	0.65

FAST and ORB, being derived from the same algorithm, are the two yielding the best performance, while at the same time finding a very high number of features. Something to note is that ORB returns approximately 8000 points per frame after removing the feature cap, and keeps frames per second in line with FAST.

Good features to track, on the other hand, offers still very good performance but returns only 25 points each frame, though being high-quality features.

Finally, at least by these metrics, SIFT seems to be the worst extractor, not finding as many points as FAST, and being very slow compared to any other algorithm.

Feature Tracking

	Features per Frame	Average FPS
GFTT + LKOF	31.08	7.41
ORB + BFMatcher	1000.0	3.53
ORB + Flann Matcher	1000.0	3.42
SIFT + BFMatcher	3471.45	0.56
SIFT + Flann Matcher	3557.33	0.49

Using Good features to track, in conjunction with Optical flow, seems to be the most performant technique.

On the other hand, on the matter of matchers, it seems that the major drop in performance is caused by the extractor's algorithm type, as there is very little difference between the average FPS, when using brute force or Flann-based matcher.

Something to note is that the performance of the optical filter is higher than just Good Features to Track alone because in this test the extractor is called less often, so the tracking

algorithm does most of the work. In the test with just point retrieval, the extractor function is called every frame.

V. TEMPLATE MATCHING EXPERIMENTS

To conclude, some experiments have been done with OpenCV's template matcher, although not very effectively. The matcher seems to find some counters but returns an offset position. My guess is that this is likely due to the large amount of similar objects in the scene, or to the uniformity in color of the video.

VI. LINKS

- [Github](#)
- [Short Output Clips](#)
- [Video Presentation](#)